

## 1 Implementation Overview

This assignment implements a parallelized Bitcoin proof-of-work miner using CUDA. The miner searches for a nonce value that, when hashed with the block header using double-SHA256, produces a hash below a target difficulty threshold.

### Architecture

- **Host (CPU):** Block parsing, merkle root calculation, result formatting
- **Device (GPU):** Parallel nonce search using optimized SHA256

## 2 Parallelization and Optimization Techniques

### 2.1 SHA256 Midstate Pre-computation

Bitcoin block headers are 80 bytes long, but only the final 4 bytes (nonce) change. SHA256 operates on 64-byte chunks, meaning the first 64 bytes can be pre-hashed once on CPU. The resulting midstate is reused for every GPU thread.

```
1 sha256_compute_midstate(first64, midstate);
2 sha256_finalize_from_midstate(&ctx1, d_midstate, last16);
```

Benefits include an 80% reduction in hash work and massive reuse of CPU precomputation.

### 2.2 Circular W-Array in SHA256

The standard SHA256 algorithm uses a 64-word schedule array, which consumes large per-thread memory. Replacing it with a 16-entry circular buffer reduces local memory usage by 75%.

```
1 WORD w[16];
2 w[i & 15] = w_i_16 + s0 + w_i_7 + s1;
```

### 2.3 Specialized sha256\_32bytes()

Double-SHA256's second hash always processes exactly 32 bytes. A custom version removes general-purpose overhead and branches.

```
1 block[32] = 0x80;
2 memset(block + 33, 0, 23);
3 block[62] = 1; block[63] = 0;
```

### 2.4 Register Caching

Repeated reads of global memory slow down SHA256. Caching values into registers eliminates repeated global memory loads.

## 2.5 Batched Atomic Operations

Instead of checking the global flag every iteration, threads check every 256 iterations, reducing contention massively.

```
1 if ((local_batch & 0xFF) == 0 && d_found) return;
```

## 2.6 Maximum Thread Utilization

Using 512 threads/block and 65535 blocks maximizes GPU parallelism without register spilling.

# 3 Experiments

## Methodology

- Test platform: Tesla V100-SXM2-32GB
- Test case: public testcase 01
- Metric: Kernel execution time (1 run)

Config	Blocks	Threads	Total Threads	Time (s)
1	256	128	32768	15.36
2	512	128	65536	10.33
3	1024	256	262144	9.64
4	4096	256	1,048,576	5.07
5	16384	256	4,194,304	5.46
6	32768	256	8,388,608	6.03
7	32768	512	16,777,216	3.56
8	65535	512	33,553,920	3.58

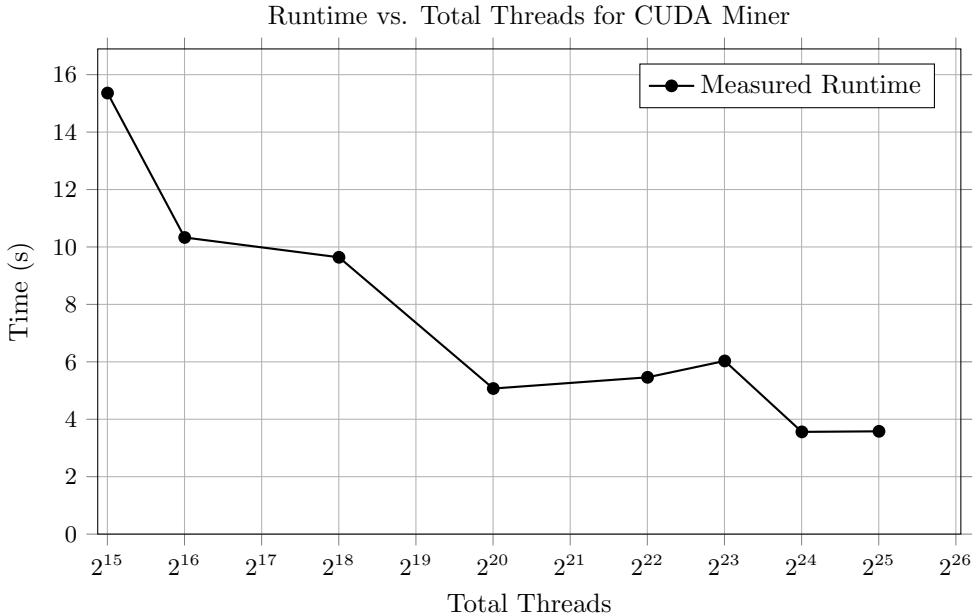


Figure 1: Runtime of CUDA Bitcoin Miner under different block/thread configurations. The x-axis is plotted in log scale to reveal scaling trends.

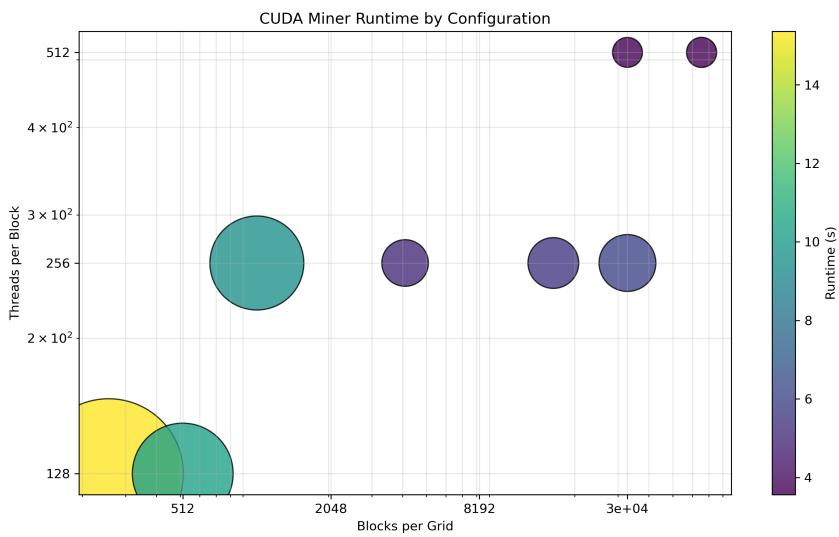


Figure 2: CUDA Miner Runtime for Different Configurations. Bubble size represents relative runtime, color corresponds to runtime in seconds. X-axis: Blocks per Grid, Y-axis: Threads per Block (both powers-of-2).

## 4 Other Details of CUDA programming

### 4.1 Constant Memory

```
1 __constant__ unsigned char d_target[32];
2 __constant__ WORD d_midstate[8];
```

### 4.2 Grid-Stride Loops

```
1 for (unsigned long long nonce = tid;
2     nonce < max_nonce_space;
3     nonce += total_threads) {
4     ...
5 }
```

### 4.3 Atomic Compare-And-Swap

```
1 if (atomicCAS(&d_found, 0, 1) == 0) {
2     d_solution_nonce = (unsigned int) nonce;
3 }
```